Prediction models for Progression-Free Survival in Breast Cancer Patients

Shivalika Chavan (sc5826), Kaveri Uberoy(kiu2103), Khue Nguyen (nnk2114)

2026-03-25

1. Introduction

Two sets of data about the progression-free survival of breast cancer patients were provided, originating from two similar studies. Dataset A comprises data from 2000 patients collected between 1990 and 1995. Dataset B was subsequently collected from 2000 to 2005, utilizing an identical study design, sample size, and set of variables.We built a prediction model for log-transformed PFS using the dataset A and evaluated the model's generalization of this model to dataset B.

2. Method

We cleaned both datasets and factorized the categorical variables by their specified levels including: menopause, race, tumor_grade, stage, surgery, chemo, hormonal_therapy, ER, PR and HER2. We checked for NAs in the data and removed the id variable as it was not predictive. We explored the correlation of variables to determine multicollinearity by checking the correlation heatmap . We set our seed as 1, and then randomly partitioned the dataset A into training (80%) and testing data (20%) using initial_split() and used repeated (n=5) 10-fold cross-validation to estimate out-of-sample prediction error. This approach repeatedly partitions the training data into training and validation sets, reducing variance compared to a single split. Cross-validation folds were explicitly created using createFolds() with k = 10 and reused across all models.

We evaluated seven candidate models: Linear Regression, K-Nearest Neighbors, Elastic Net

(ENET), Principal Component Regression (PCR), Partial Least Squares (PLS), Generalized Additive Models (GAM), and Multivariate Adaptive Regression Splines (MARS). For KNN, the tuning parameter k was evaluated over the range 1 to 25, and all predictors were pre-processed by centering and scaling prior to model fitting. For ENET, the mixing parameter alpha was tuned over a grid defined by seq(0,1,length=11), and the regularization parameter lambda was tuned over a grid defined by exp(seq(−5,2,length=50)). For MARS, the degree of interactions was tuned over the range 1 to 3, and the number of terms retained in the model (nprune) was tuned over the range 2 to 15.For GAM, the default tuning parameters provided by the caret package were used. For Principal Component Regression (PCR) and Partial Least Squares (PLS), the optimal number of components was selected using a tuning length of 20 within the cross-validation framework, and all predictors were pre-processed by centering and scaling prior to model fitting.

Modeling was done using R version 4.5.1. Packages included "tidymodels" for cleaning data, "car" for assessing variance inflation factors (VIF), "caret" and "splines" for model training, "pls" for partial least squares model, "earth" for MARS model, "patchwork" for plot layout.

3. Results and discussion

3.1. Exploratory Data Analysis

Both dataset A and dataset B each contained data of 2000 patients with complete follow-up and no missing values. There are 15 predictors, including 5 continuous and ten categorical. The sample characteristics are described in Table 1.

Figure 1 and 2 show the correlation between variables in dataset A and B respectively . First, to look at data A, there are a few correlated variables, which intuitively make sense:

- `tumor_size`, `lymph_nodes`, and `stage` are highly positively correlated
- `age` and `menopause` are positively correlated
- `Ki67` (measures how fast cancer cells in a tumor are dividing) and `tumor_grade` (how abnormal cancer cells look) are positively correlated
- `Ki67` and `ER` (estrogen receptor) are negatively correlated

- `PR` (progesterone receptor) and `ER` are positively correlated
- `hormonal_therapy` and `ER` are positively correlated

We checked variance inflation factors (VIF) values to assess the multicollinearity. All adjusted GVIF values were less than 2.5, indicating no multicollinearity. The highest value was observed for `tumor_size` (2.17), suggesting moderate but acceptable correlation with other variables (`stage` and `lymph node`), see Table 2. Thus, all variables were included in the modelling.

3.2. Model selection for dataset A

The cross-validated RMSE for each model was computed as the mean RMSE across all resampling iterations. Based on the training data, most models performed very similarly, with mean RMSE values ranging from 0.7054 (PLS) to 0.7542 (GAM). (see Table 3). The similar performance across linear and nonlinear models suggests that the relationship between the predictors and `log(PFS)` is largely linear. PLS achieved the lowest mean error (0.7054), closely followed by PCR (Mean = 0.7066) and ENET (Mean = 0.709). ENET and PCR performing similarly to PLS suggests that dimensionality reduction provided little additional predictive power. ENET has the lowest maximum RMSE observed. Except for one outlier, the remaining RMSEs from cross-validation have a tight distribution, indicating this model is less variable and more stable.

Because predictive accuracy was comparable across the top models, we prioritized interpretability. We selected the ENET model, which was trained using an explicit lambda grid to ensure consistent tuning across folds. The cross-validation process selected an alpha of 0 (with lamba of 0.751), meaning the final model is a pure Ridge regression. While Ridge regression does not perform sparse feature selection (all variables are retained), it stabilizes the coefficients to handle multicollinearity, making it a highly robust and easily interpretable model for our clinical context. The coefficients for each variable of our selected model are shown in Table 4. The reported test error for this ENET model was 0.7445; test RMSE was calculated by comparing predicted and observed values of `log(PFS)` using the root mean squared error metric.

3.3. Model testing and new model for dataset B

3.3.1. Model testing on dataset B

When applying the final Elastic Net model (trained on datA) to the new cohort (datB), the RMSE increased to 0.8422. Compared to the previous test RMSE of 0.7445 on the original data, this indicates a drop in predictive performance. This gap suggests that there may be underlying distribution shifts between the two patient populations (e.g., differences in demographics, clinical settings, or unmeasured variables), or that the model slightly overfitted the original cohort despite the use of Ridge regularization. A new model will be needed.

3.3.2. Train new model on dataset B

Since ENET had a tight distribution when using the original data, it seems the most promising (rather than training many different models). Additionally, it is then easier to compare what changed in the model. Like the old model trained on dataset A, the best tuning parameters remained similar ($\alpha = 0$, $\lambda = 0.736$).

By comparing the coefficients of the old pre-2000 model to the new 2000-2005 model (see Table 5), several major predictive shifts become clear. The new model's intercept increased from 2.55 to 3.03, indicating that the baseline PFS increased over the decade. This makes sense since new and better interventions will have likely been more accessible. The impact of HER2 status (`HER21`) has flipped. In the older cohort, being HER2+ was a severe negative predictor for progression-free survival (-0.101). In the new cohort, this penalty essentially vanished (+0.009, which is close to 0). Other predictors that flipped were all race categories, `menopause1`, `stage2`, and `tumor_grade2`. Other variables including ER and hormonal therapy stayed largely the same and continued to improve PFS, while stage 3, tumor grade 3, and tumor size decreased PFS. The gap in predictive power makes sense given these changes in cohort trends (see Table 5).

Appendix

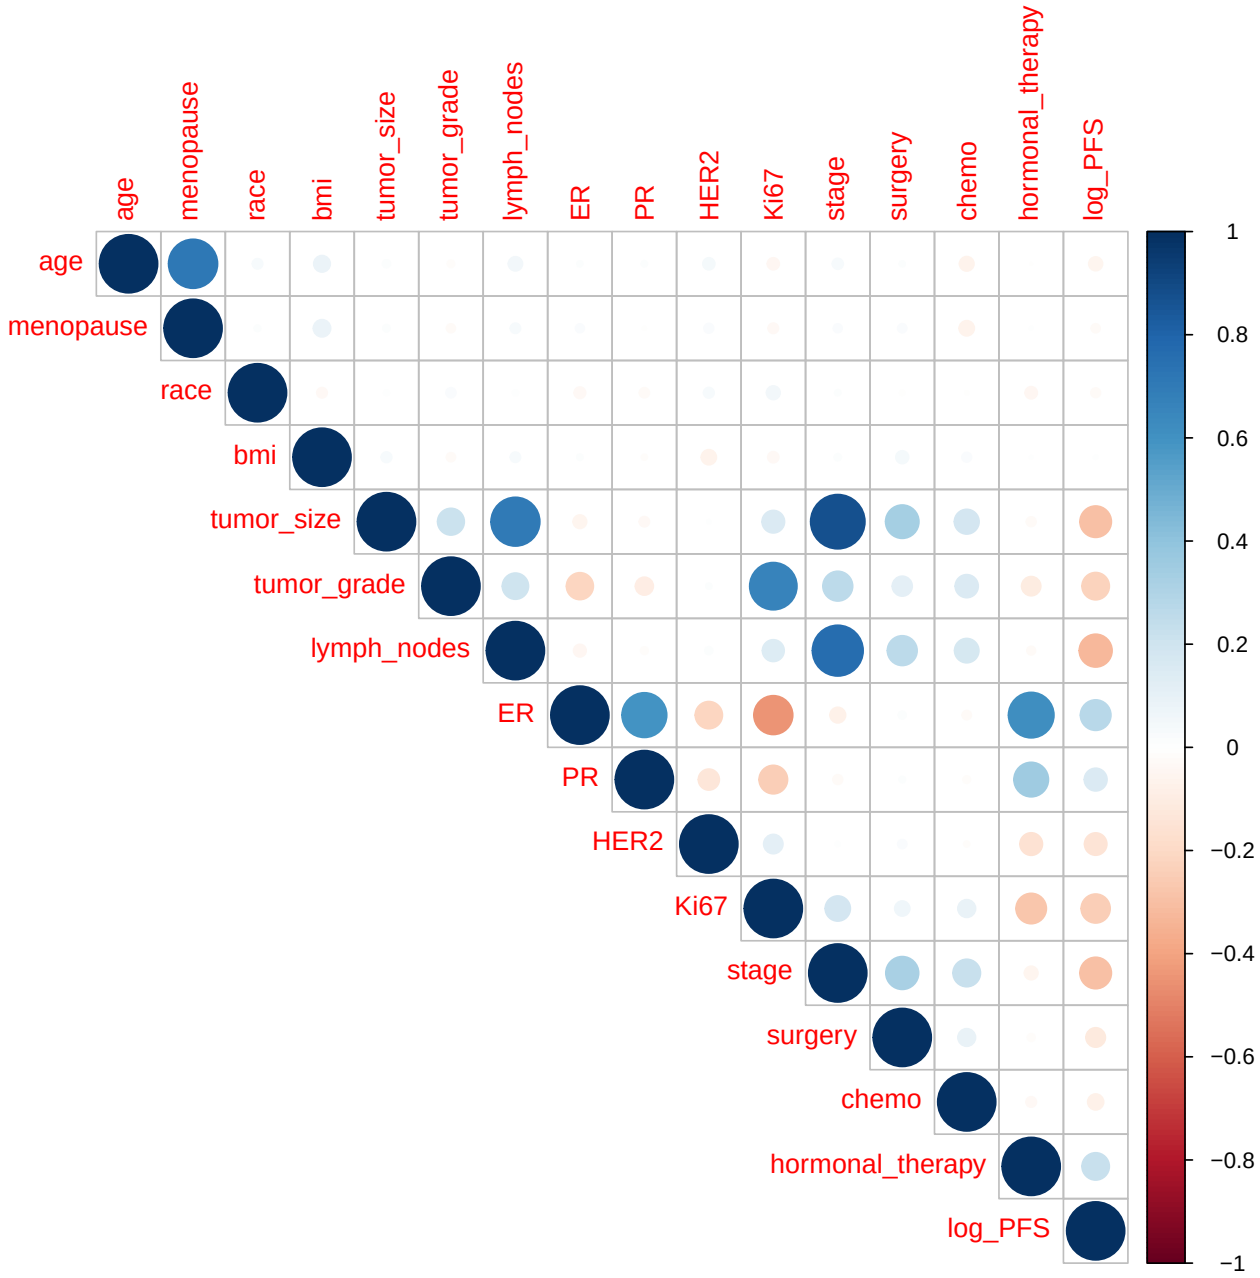

Figure 1: Correlation matrix of variables for dataset A

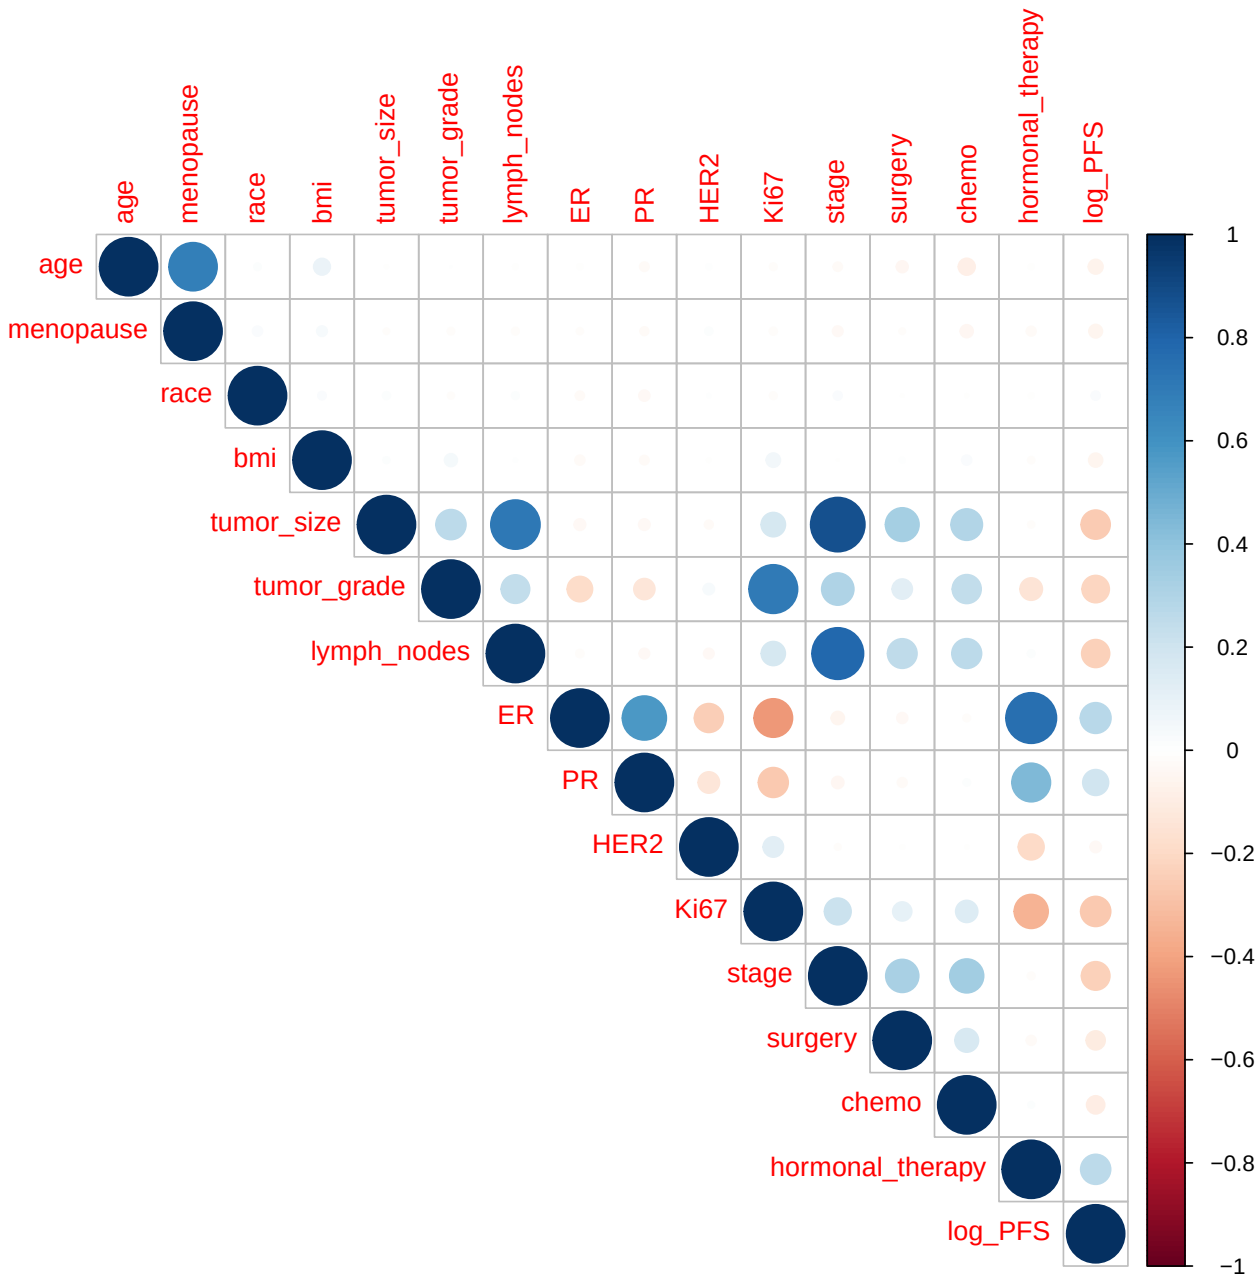

Figure 2: Correlation matrix of variables for dataset B

Table 1: Sample characteristics of Dataset A and Dataset B.

Characteristic	**N = 2,000**	**N = 2,000**
age	55 (47, 62)	56 (49, 64)
menopause	NA	NA
0	689 (34%)	638 (32%)
1	1,311 (66%)	1,362 (68%)
race	NA	NA
1	1,284 (64%)	1,097 (55%)
2	289 (14%)	369 (18%)
3	176 (8.8%)	220 (11%)
4	251 (13%)	314 (16%)
bmi	26.1 (23.0, 29.1)	27.3 (24.0, 30.9)
tumor_size	3.10 (2.10, 4.80)	2.80 (1.80, 4.20)
tumor_grade	NA	NA
1	383 (19%)	502 (25%)
2	871 (44%)	839 (42%)
3	746 (37%)	659 (33%)
lymph_nodes	2 (1, 4)	2 (0, 4)
ER	NA	NA
0	719 (36%)	652 (33%)
1	1,281 (64%)	1,348 (67%)
PR	NA	NA
0	1,039 (52%)	991 (50%)
1	961 (48%)	1,009 (50%)
HER2	NA	NA
0	1,556 (78%)	1,526 (76%)
1	444 (22%)	474 (24%)

Table 1: Sample characteristics of Dataset A and Dataset B. *(continued)*

Characteristic	**N = 2,000**	**N = 2,000**
Ki67	22 (14, 30)	20 (12, 29)
stage	NA	NA
1	504 (25%)	692 (35%)
2	893 (45%)	828 (41%)
3	603 (30%)	480 (24%)
surgery	NA	NA
0	986 (49%)	1,117 (56%)
1	1,014 (51%)	883 (44%)
chemo	NA	NA
0	895 (45%)	702 (35%)
1	1,105 (55%)	1,298 (65%)
hormonal_therapy	NA	NA
0	1,066 (53%)	861 (43%)
1	934 (47%)	1,139 (57%)
log_PFS	2.19 (1.67, 2.71)	2.74 (2.25, 3.25)

Table 2: Variance Inflation Factors (VIF)

Variable	GVIF	Df	GVIF^(1/(2*Df))
age	2.120	1	1.456
menopause	2.115	1	1.454
race	1.031	3	1.005
bmi	1.020	1	1.010
tumor_size	4.712	1	2.171
tumor_grade	1.988	2	1.187
lymph_nodes	2.655	1	1.630
ER	2.538	1	1.593
PR	1.537	1	1.240
HER2	1.066	1	1.032
Ki67	2.192	1	1.481
stage	6.316	2	1.585
surgery	1.128	1	1.062
chemo	1.077	1	1.038
hormonal_therapy	1.635	1	1.279

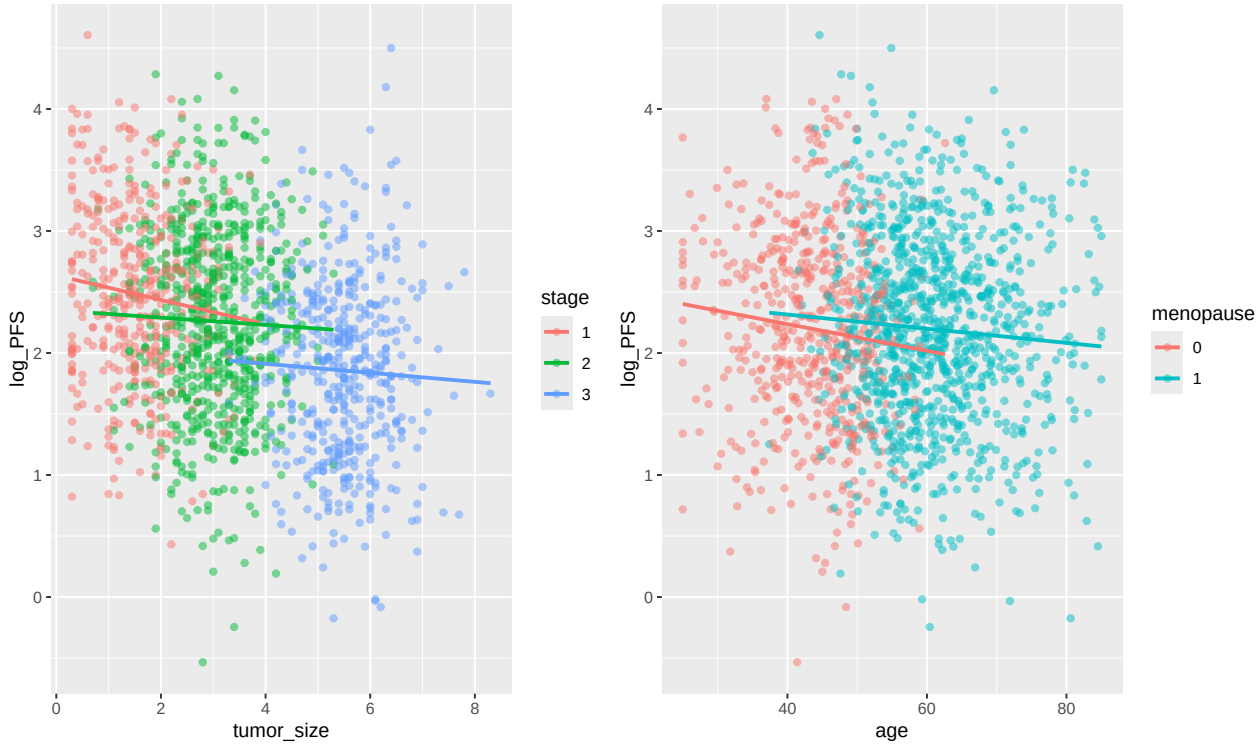

Figure 3: Interaction plot between tumor_size and stage, menopause and age

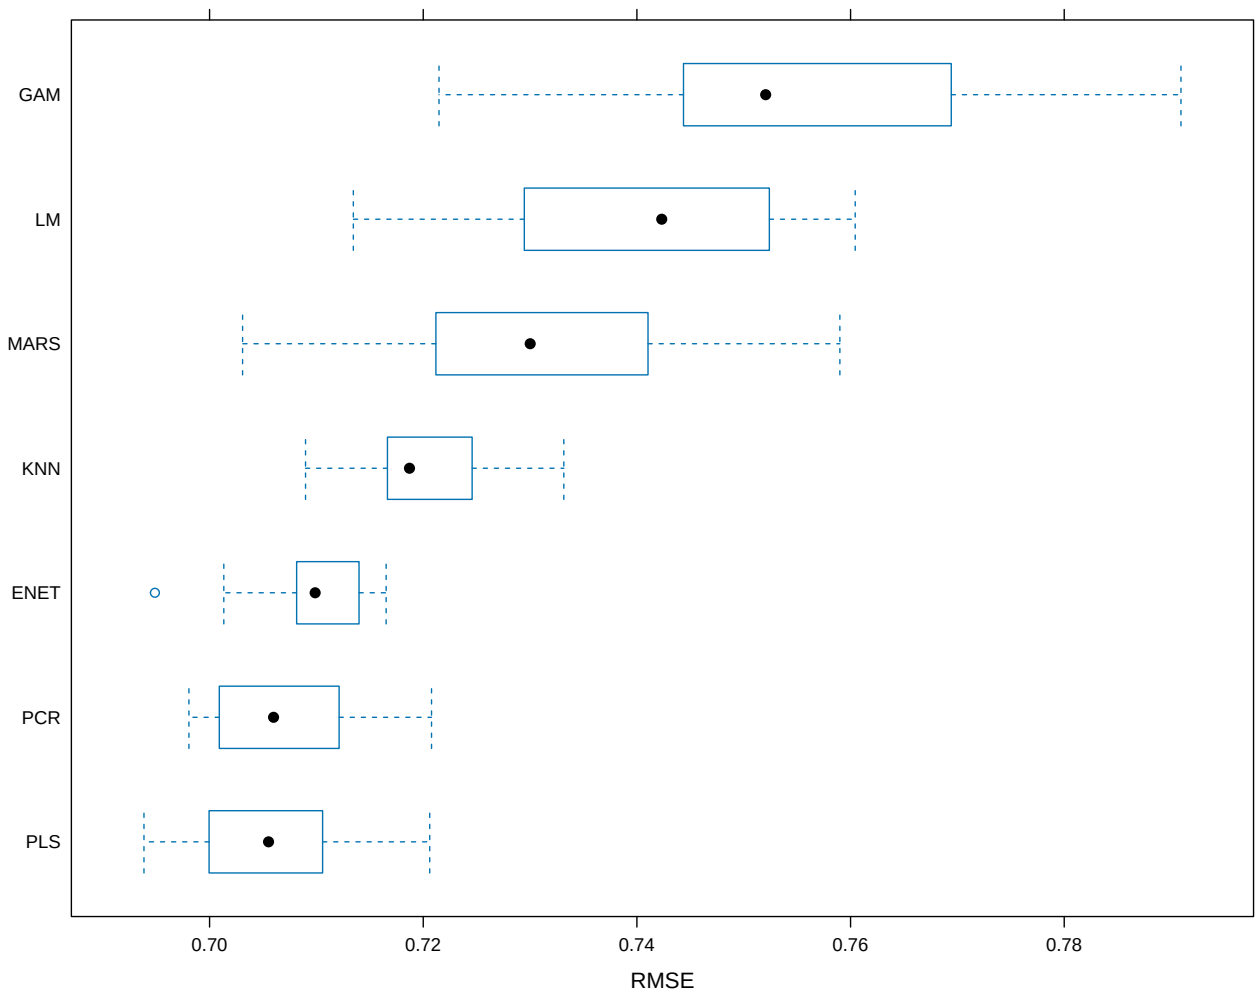

Figure 4: RMSE of tested models

Table 3: RMSE comparison for tested models.

model	cv_rmse
LM	0.740
KNN	0.720
ENET	0.709
PCR	0.707
PLS	0.705
MARS	0.754
GAM	0.730

Table 4: Elastic Net Coefficients for Dataset A.

predictor	coefficient
(Intercept)	2.550
ER1	0.129
hormonal_therapy1	0.106
stage3	-0.102
HER21	-0.101
tumor_grade3	-0.088
PR1	0.038
tumor_size	-0.034
surgery1	-0.033
race4	-0.032
lymph_nodes	-0.030
race3	0.028
race2	0.019
menopause1	0.015
chemo1	-0.012
tumor_grade2	-0.006
Ki67	-0.004
age	-0.002
bmi	0.000
stage2	0.000

Table 5: Coefficients comparision of two ENET models.

predictor	coefficient_datasetA	coefficient_datasetB	difference
(Intercept)	2.5499	3.0341	-0.4842
ER1	0.1286	0.1117	0.0169
HER21	-0.1009	0.0086	-0.1096
Ki67	-0.0042	-0.0041	-0.0001
PR1	0.0384	0.0611	-0.0227
age	-0.0021	-0.0019	-0.0002
bmi	-0.0002	-0.0031	0.0028
chemo1	-0.0119	-0.0147	0.0027
hormonal_therapy1	0.1058	0.1224	-0.0166
lymph_nodes	-0.0300	-0.0158	-0.0142
menopause1	0.0149	-0.0220	0.0369
race2	0.0195	-0.0119	0.0314
race3	0.0284	-0.0284	0.0567
race4	-0.0318	0.0478	-0.0796
stage2	0.0000	-0.0105	0.0105
stage3	-0.1017	-0.0842	-0.0176
surgery1	-0.0325	-0.0476	0.0150
tumor_grade2	-0.0065	0.0568	-0.0633
tumor_grade3	-0.0881	-0.0822	-0.0059
tumor_size	-0.0343	-0.0293	-0.0050